

Michael Lubinsky

mlubinsky@hotmail.com mobile: 408 775 4518 Nationality: USA

Summary

More than 10 years experience building data pipelines and doing data analysis.

Technical Skills

- Big Data: Apache Spark, Apache Flink, Hadoop, Snowflake
- Databases: PostgreSQL, MySQL, MariaDB, Oracle, Redshift, Presto, Druid, Hana, DynamoDB
- Languages: Python, Java, Scala, C, C++
- Apache Airflow, Kafka, Apache NiFi
- Cloud platforms: AWS, GCP, Databricks
- Data analysis and machine learning: Pandas, sklearn, PyTorch, Tensorflow
- Reporting tools: Looker, Tableau, Superset

Working Experience

Roku, San Jose, CA Apr 2020 – present. Senior Data Engineer

- Developing the data pipelines using Apache Airflow, Apache Spark, Redshift, Snowflake and Kafka
- Data acquisition, database design and implementation, dashboard/reporting using Looker and Tableau
- Working with different data formats: JSON, Parquet, Avro
- Deploying ML models to production, MLOps, Data governance

ARM, San Jose, CA Apr 2018 – 2020. Senior Software Engineer

- Developed the data pipelines for IoT (Internet of Things): data acquisition using MQTT protocol
- Created database schemas, tables, stored procedures, materialized views
- Implemented time series data analysis and anomaly detection algorithms
- Environment and tools: AWS S3, Hadoop, Hive, Python, Java, SQL, Docker

Apple, Cupertino, CA Nov 2013 – Mar 2018. Senior Software Engineer

- Developed the high performance, scalable and distributed pipeline: data acquisition, parsing, uploading, analysis and reporting.
- Implemented data clustering, supervised and non-supervised learning algorithms.
- Developed database models and REST API, which is used by external applications.
- Environment and tools: Java, C++, Python, PostgreSQL and no-SQL databases, RabbitMQ, Hadoop, Tableau.

SAP, Palo Alto, CA Nov 2011 – Oct 2013. Senior Software Engineer

- Created the logical and physical data modeling for HANA database; data partitioning and replication.
- Database performance profiling and tuning: # of threads, memory allocation.
- Run TPC-H database benchmarking
- Environment and tools: Linux, C++; Java, Python and SQL.

Yahoo Inc., Sunnyvale, CA Nov 2010 – Nov 2011. Senior Software Consultant in User Data and Analytics Group

- I wrote map-reduce jobs for processing the hundreds of GB of data on Hadoop using Java and Pig Latin.
- I implemented the sampling algorithm, which dynamically allocates the number of Hadoop reducers per key proportionally to the data distribution
- Applied statistical methods (regression, classification) for Apache log data analysis.
- Environment and tools: Linux, Java, Hadoop.

Education & Qualifications

- Apache Spark, Udemy 2015
- Introduction to Artificial Intelligence, Coursera 2011
- Machine Learning, Coursera 2011
- MS in EE (with honor) - Moscow Institute of Electronic Engineering